

Image Segmentation using Neural Networks


Jakub Náplava, Jan Klůj, Ondřej Švec

What is Segmentation?

- semantic segmentation
- instance segmentation

<http://stackoverflow.com/questions/33947823/what-is-semantic-segmentation-compared-to-segmentation-and-scene-labeling>

http://www.cs.toronto.edu/~urtasun/courses/CSC2541/08_instance.pdf



Architectures

- FCN (Nov 14)

https://people.eecs.berkeley.edu/~jonlong/long_shelhamer_fcn.pdf

- SegNet (Nov 15)

<https://arxiv.org/pdf/1511.00561.pdf>

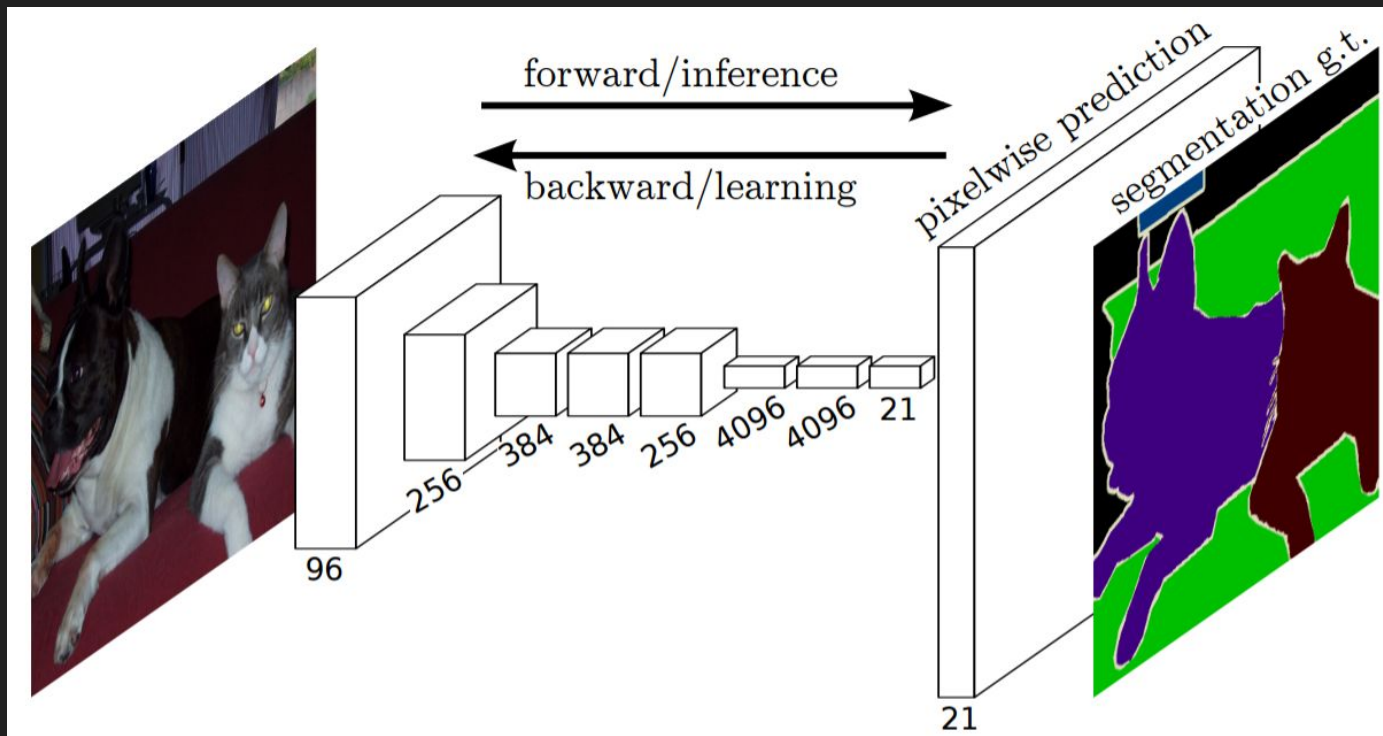
- DeconvNet (May 15)

<https://arxiv.org/pdf/1505.04366v1.pdf>

- DeepLab-LargeFOV (Dec 14)

<https://arxiv.org/pdf/1412.7062v4.pdf>

FCN Fully Convolutional Networks for Semantic Segmentation



FCN Fully Convolutional Networks for Semantic Segmentation

- VGG16/GoogLeNet architectures
- conv layer 1x1, #channels = #classes
- one upsampling layer (transposed deconv)
- skip connections





Figure 3. Our DAG nets learn to combine coarse, high layer information with fine, low layer information. Pooling and prediction layers are shown as grids that reveal relative spatial coarseness, while intermediate layers are shown as vertical lines. First row (FCN-32s): Our single-stream net, described in Section 4.1, upsamples stride 32 predictions back to pixels in a single step. Second row (FCN-16s): Combining predictions from both the final layer and the pool4 layer, at stride 16, lets our net predict finer details, while retaining high-level semantic information. Third row (FCN-8s): Additional predictions from pool3, at stride 8, provide further precision.

SegNet

A Deep Convolutional Encoder-Decoder Architecture for Image Segmentation




SegNet

A Deep Convolutional Encoder-Decoder Architecture for Image Segmentation


- VGG16 architecture
- special demaxpool
- conv layers after demaxpool

Convolution with trainable decoder filters



SegNet

Deconvolution
for upsampling



Encoder feature map
FCN

DeconvNet

Learning Deconvolution Network for Semantic Segmentation


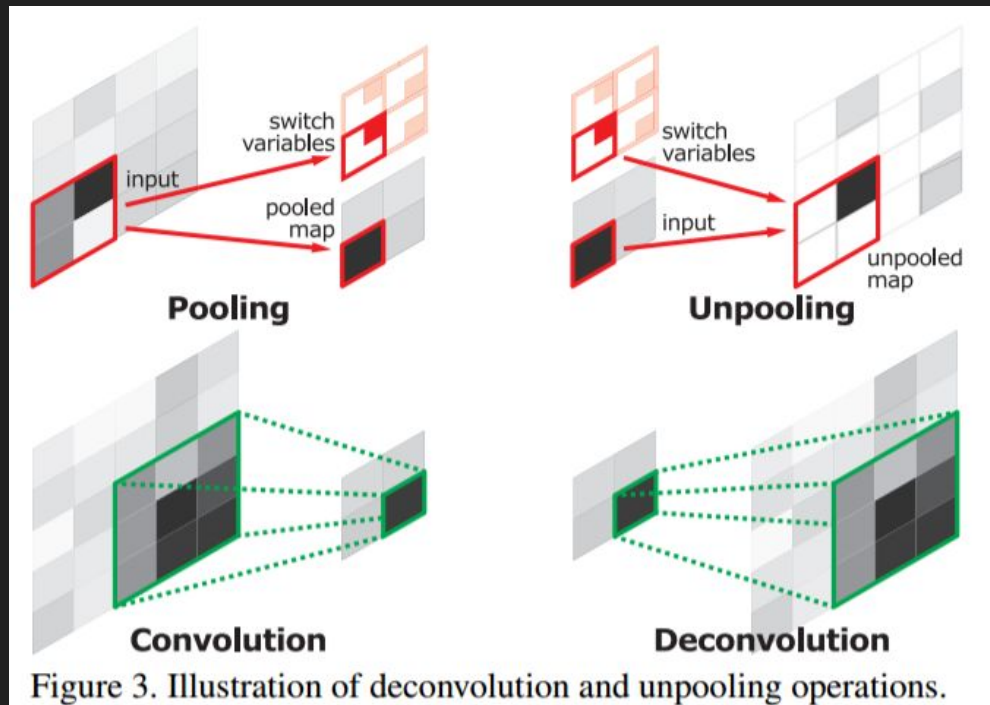


Figure 2. Overall architecture of the proposed network. On top of the convolution network based on VGG 16-layer net, we put a multi-layer deconvolution network to generate the accurate segmentation map of an input proposal. Given a feature representation obtained from the convolution network, dense pixel-wise class prediction map is constructed through multiple series of unpooling, deconvolution and rectification operations.


Upsampling methods

- bilinear interpolation
- transposed convolution
- naive demaxpool
- learnable demaxpool
- fixed demaxpool



Dataset

- http://vis-www.cs.umass.edu/lfw/part_labels/
- part of 'Labeled Faces in the Wild' dataset
- 2927 samples




Data augmentation

- crop
- flip
- gaussian blur
- dropout
- additive gaussian noise
- affine


GaussianBlur



AdditiveGaussianNoise




Affine: Translate




Dropout



Crop
(top, right,
bottom, left)



Fliplr



Common accuracy metrics

- per pixel accuracy
- mean accuracy
- mean intersection over union (IU)
- frequency weighted IU

$$\sum_i n_{ii} / \sum_i t_i$$

$$(1/n_{cl}) \sum_i n_{ii} / t_i$$


$$(1/n_{cl}) \sum_i n_{ii} / \left(t_i + \sum_j n_{ji} - n_{ii} \right)$$

$$\left(\sum_k t_k \right)^{-1} \sum_i t_i n_{ii} / \left(t_i + \sum_j n_{ji} - n_{ii} \right)$$

n_{ij} be the number of pixels of class i predicted to belong to class j , where there are n_{cl} different classes, and let $t_i = \sum_j n_{ij}$ be the total number of pixels of class i .

$$t_i = \sum_j n_{ij}$$

Our architecture




Experiments and results

- Architecture: no max pools, no shared weights, no skip connections

⇒ 92.92% (per pixel accuracy)

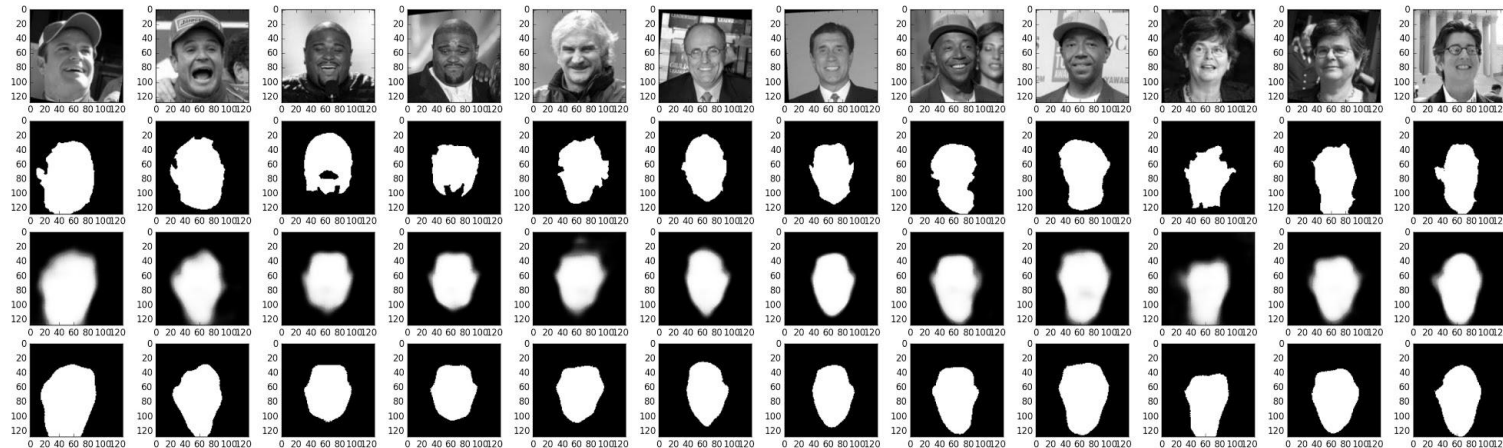
Accuracy: 0.9292570948600769, Filters: [64, 64, 64, 64, 128, 128, 128, 128], Sizes: [3, 3, 3, 3, 7, 7, 7, 7]



Experiments and results

- Architecture: No shared weights, no skip connections \Rightarrow 94.62%
 - 4 conv + 4 deconv

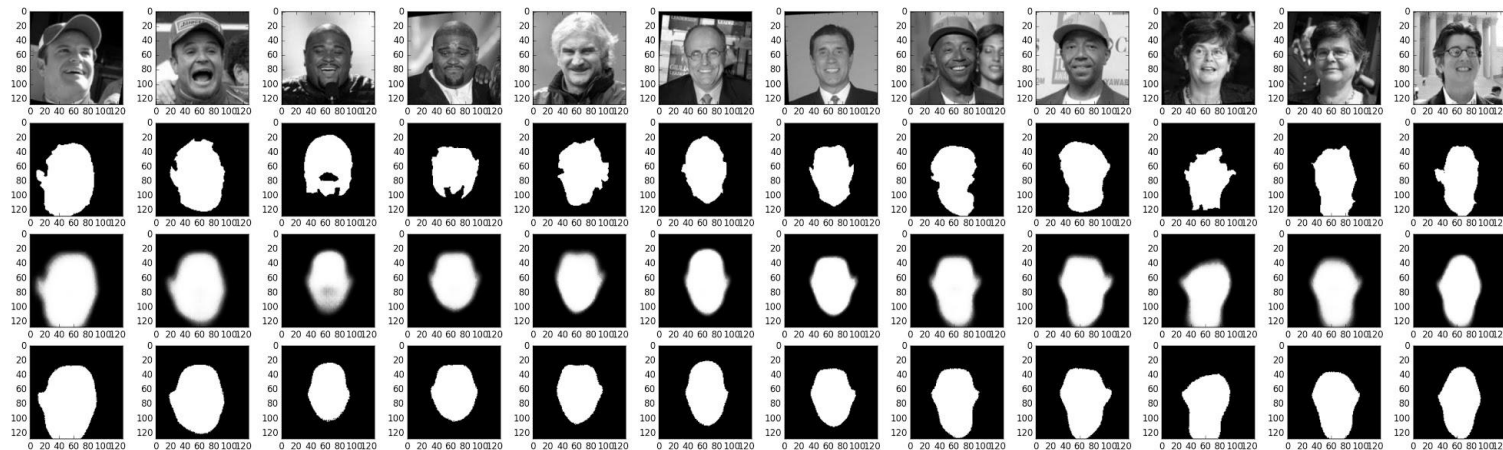
Accuracy: 0.9462901949882507, Net: C:64K:7, C:64K:7, MP:2, C:64K:7, C:64K:7, MP:2,



Experiments and results

- Architecture: more layers... 8 conv + 8 deconv \Rightarrow 96.14%

Accuracy: 0.9614855647087097, C7,64,2C7,64,1M2C7,64,2C7,64,1M2C7,128,2C7,128,1M2C7,128,2C7,128,1



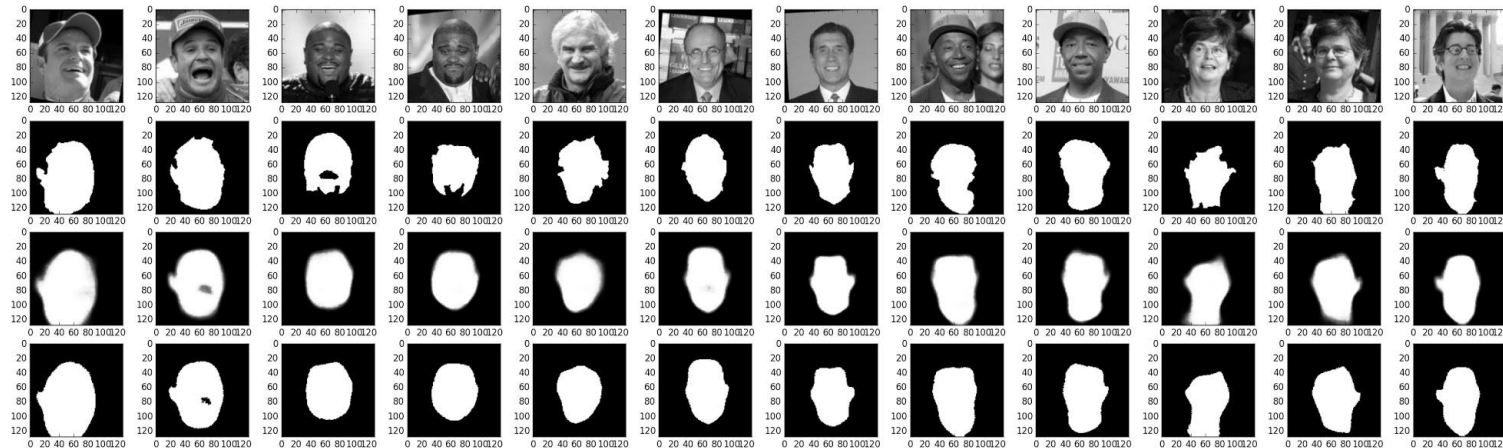
Experiments and results

- Architecture:

- kernels - 7x7
- channels - 64
- shared weights between conv & deconv

96.64 %

Accuracy: 0.9664298295974731, C7,64,2C7,64,1M2C7,64,2C7,64,1M2C7,64,2C7,64,1M2



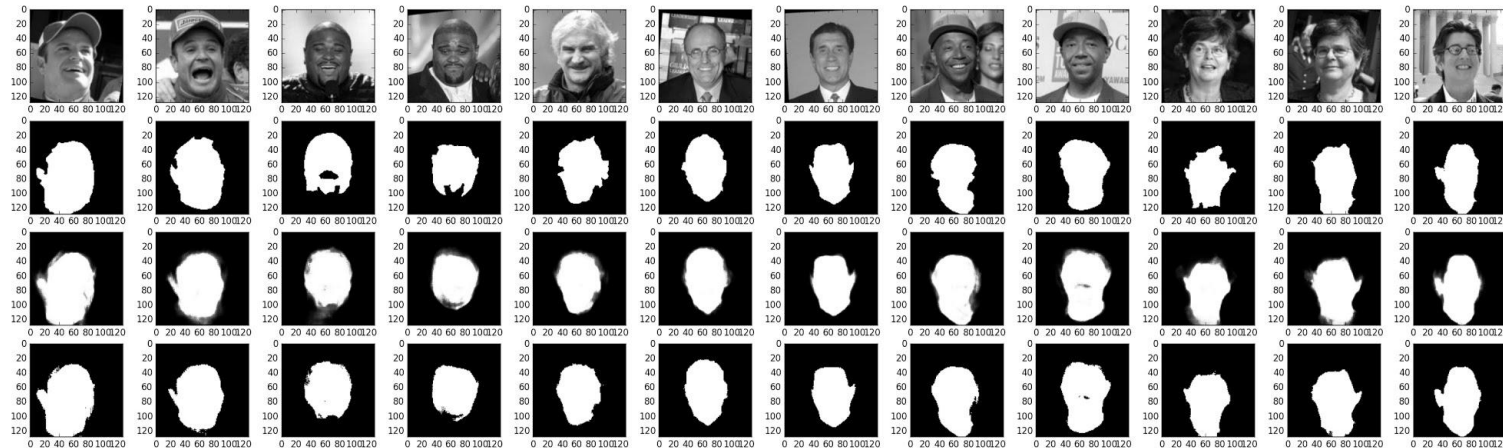
Experiments and results

- Architecture:

- kernels - 7x7
- channels - 64
- shared weights between conv & deconv
- added skip connections

97.21 %

Accuracy: 0.9721238017082214, C7,64,2C7,64,1M2C7,64,2C7,64,1M2C7,64,2C7,64,1M2

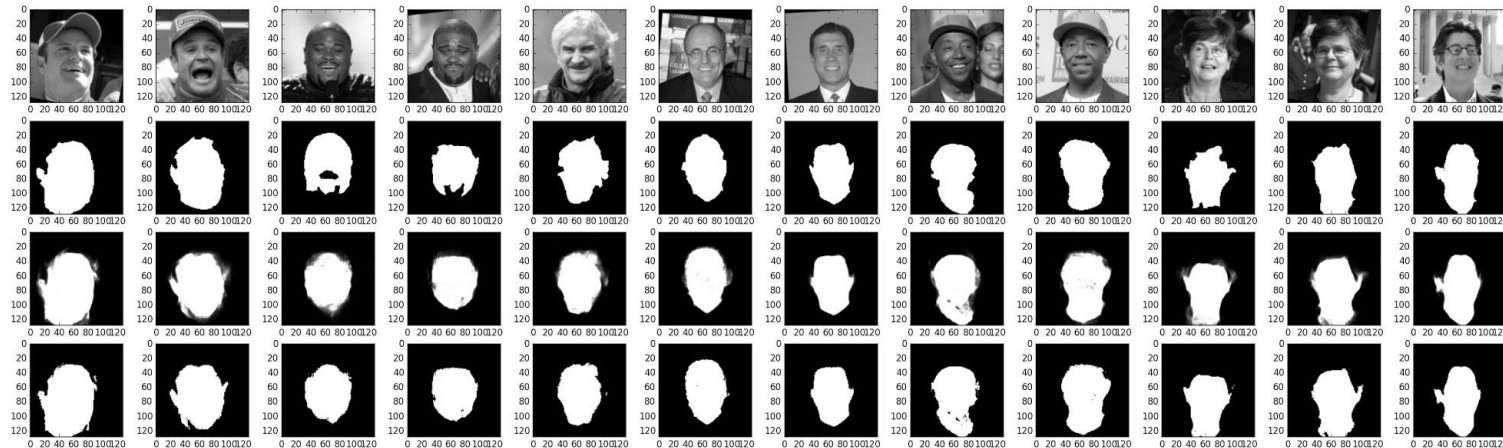


Experiments and results

- Architecture:
 - kernels - 7x7
 - channels - 64
 - shared weights between conv & deconv
 - added skip connections + image standardization

97.22 %

Accuracy: 0.9722825884819031, C7,64,2C7,64,1M2C7,64,2C7,64,1M2C7,64,2C7,64,1M2



Experiments and results

- The best...
- kernels - 5x5

97.36 %

Accuracy: 0.9736009240150452, C5,64,2C5,64,1M2C5,64,2C5,64,1M2C5,64,2C5,64,1M2



What did not help...

- CPU
- All strides = 1 (in conv layers)
- Batch-norm

Thank you!

Questions?

